

```
// VDotProduct.cu
//
// driver and kernel call

#include <stdio.h>
#include <cuda.h>
#include <stdlib.h>

#define THREADS_PER_BLOCK 32

__global__ void dot_product (int *a_d, int *b_d, int *c_d, int arraySize)
{
    int x = blockIdx.x * blockDim.x + threadIdx.x;
    int t = threadIdx.x;
    __shared__ int temp[THREADS_PER_BLOCK];

    if (x < arraySize/2)
    {
        a_d[x] = x + 1;
    }
    else
    {
        a_d[x] = x - ((x - arraySize/2)*2);
    }

    b_d[x] = (x % 10) + 1;

    temp[t] = a_d[x] * b_d[x];
    __syncthreads();

    if(0 == threadIdx.x)
    {
        int sum=0;
        for(int i=0;i<THREADS_PER_BLOCK;i++)
        {
            sum+=temp[i];
        }
        atomicAdd(c_d,sum);
    }
}

extern "C" void gpu_dot_product (int *a, int *b, int *c, int arraySize)
{
    int *a_d, *b_d, *c_d;

    cudaMalloc ((void**) &a_d, sizeof(int) * arraySize);
    cudaMalloc ((void**) &b_d, sizeof(int) * arraySize);
    cudaMalloc ((void**) &c_d, sizeof(int) * arraySize);
    cudaMemcpy (a_d, a, sizeof(int) * arraySize, cudaMemcpyHostToDevice);
    cudaMemcpy (b_d, b, sizeof(int) * arraySize, cudaMemcpyHostToDevice);
```

```
dot_product <<< ceil((float) arraySize/THREADS_PER_BLOCK), THREADS_PER_BLOCK >>> (a_d, b_d,
c_d, arraySize);

cudaError_t err = cudaGetLastError();
if (err != cudaSuccess)
    printf ("CUDA error: %s\n", cudaGetErrorString(err));

cudaMemcpy (c, c_d, sizeof(int) * arraySize, cudaMemcpyDeviceToHost);
cudaFree (a_d);
cudaFree (b_d);
cudaFree (c_d);

}
```